

# SY19 - Apprentissage à partir de trois jeux de données

*Jingyi HU*

*Décembre 6, 2017*

The goal of this TP is to build classifiers as powerful as possible from three real data sets. I have 3 datasets:

- `expression_train.txt`, containing the gray levels of an image associated with a facial expression among joy, surprise, sadness, disgust, anger, fear;
- `characters_train.txt`, consisting of examples from 26 classes corresponding to the 26 letters of the alphabet;
- `speech_train.txt`, where each observation corresponds to pronunciation of a phoneme among “sh”, “dcl”, “iy”, “aa” and “ao”. I will therefore compare the predictive qualities of several models on these datasets.

NB: If I have enough time, I may try a selection of main components by Forward Stepwise Selection. Because the analysis ideas of the 3 datasets are almost same, I will analyze more precisely for the first one. And I will write only several codes for the second and third dataset.

## Load of data

```
data_expressions<-read.table("expressions_train.txt")
data_characters<-read.table("characters_train.txt")
data_parole<-read.table("parole_train.txt")
```

## Using methods

1. The dimension reduction methods:

- Principal Component Analysis (PCA)
- Factor Discriminant Analysis (FDA)
- Forward Stepwise Selection
- ...

2. Comparison between the performances of these different models:

- K nearest neighbors (KNN)
- Linear Discriminant Analysis (LDA) / Quadratic Discriminant Analysis (QDA) / Regularized Discriminant Analysis (QDA) / Naive Bayesian Classifier
- Random Forest (RF)
- Support Vector Machines (SVM)
- Neural Networks (NN)

3. In addition, I also thought about the possibility to proceed by **cross validation** in order to be able to learn the models on more data, and I tried the 2 at the same time (without and with CV) I get almost the **same result**. Indeed, the number of observations available is quite low. As a result, I didn't implement the cv to methods because the results obtained seemed good.
4. I didn't try the **Logistic Regression and Generalized Additive Models (GLM & GAM)**, because this classification problem is not a binary classification. I didn't use logistic regression because the logistic regression method performed by GLM and GAM (which is an extension of GLM) is often less efficient than discriminant analysis (LDA, QDA ...) for classifications with more than two classes. In addition, the latter is often used to explain the effects of the different predictors, which is not the purpose of our study. As a result, these methods will not be used for the "word" and "characters" data sets either.

## 1.translation\_expression.txt

### First analysis

#### Pre-treatment

By displaying the different images, I also realize that part of each image is totally useless in the context of our analysis, or even risk noises: it is all the black pixels at the bottom of each image. I decide to exclude them from any future analysis.

#### "Average faces"

To better analyze this database, I tried to average for each expression of images: that is to say, produce six images from the mathematical average of the images of the expression considered.

#### Methods to use

After studying the dataset, I found that it has **a large number of predictors compared to the number of individuals**. The different predictors are globally very **correlated**. I will therefore have to use dimension reduction methods on this dataset, and then test our different models.

#### The dimension reduction methods

##### PCA

I used the PCA using the *prcomp* function of R. In order to retain only a certain number of principal components, I'm interested in the proportion of variance explained by each of them by plotting their explained proportion of variance, and cumulative variance. You can find a pdf named *pc\_cumpropvar.pdf*, *pc\_propvarexp.pdf* and *pc\_propvarexp2.pdf*. In order to test the performance of models on different selections of **Principal Component (PC)**, I composed several dataset, PCi: From the first i PCs, representing i% of explained variance. I also tried to plot our first two main component, hoping to see some interesting pattern. You can find a pdf named *pc\_2prcomp.pdf*.

## FDA

According to the poly, I used the lda method, and multiplied the matrix **X** of the individuals by the matrix *lda.scaling* of the selected eigenvectors. I applied the FDA two datasets: 1. Raw FDA: dataset containing all of the individuals in the initial dataset 2. PC100 FDA: dataset containing the 100 PCs of all individuals in the initial dataset. According to the result, I found that I have an overfitting result, the error rate is too low, and the classes are clearly too separated. The FDA on PC100 makes me realize that I didn't use the right approach regarding the FDA. Therefore, I tried using **Cross-Validation** 6-folds on FDA, and I have a rough estimate of the FDA test error rate. This time the overfitting finished.

### Conclusion:

1. the use of PCA before FDA application and simply using FDA on the raw dataset both have good results.
2. The choice of the number of PCs is important: a large number of PCs may reproduce our previous overfitting problem, while too few of them do not allow a good representativeness of the initial dataset, because they contain a too small proportion of explained variance in this case.

As FDA on the raw dataset works well, I will analyze more precisely by directly using the **FDA on the raw dataset**.

### 1.1 FDA

```
data1<-data_expressions
set.seed(2000)
#Generate the training data and test data
N<-nrow(data1)
nbTrain=floor(3/4*N)
nbTst=N-nbTrain
trainIdx<-sample(1:N, nbTrain)
train<-data1[trainIdx,]
test<-data1[-trainIdx,]
```

Now we separate the dataset into a test set and training set.

```
#Remove column full of 0
cleanData.train<-train[,colSums(abs(train[,1:ncol(train)-1])) !=0]
cleanData.test<-test[,colSums(abs(test[,1:ncol(test)-1])) !=0]
#FDA (with train data)
library("MASS")
lda_data1<- lda(y~.,data=cleanData.train)
U<-lda_data1$scaling
X<-as.matrix(cleanData.train[,1:ncol(cleanData.train)-1])
Z<-X%*%U
Z<-as.data.frame(Z)
y<-cleanData.train$y#delete useless informtion
train.FDA<-cbind(Z,y)
#Apply FDA on test data
X<-as.matrix(cleanData.test[,1:ncol(cleanData.test)-1])
Z<-X%*%U
Z<-as.data.frame(Z)
y<-cleanData.test$y
test.FDA<-cbind(Z,y)
#plot(train.FDA[,1], train.FDA[,2], col = train.FDA$y)
```

## 1.2 K nearest neighbors (KNN)

```
library(class)
train.X <- train.FDA[,1:ncol(train.FDA)-1]
train.Y <- train.FDA[,ncol(train.FDA)]
test.X <- test.FDA[,1:ncol(test.FDA)-1]
test.Y <- test.FDA[,ncol(test.FDA)]

nbvoisin <- seq(1,65)
error <- rep(1:65)
for (i in (1:65)){
  knn.pred=knn(train.X,test.X,train.Y,k=i)
  table(knn.pred,test.Y)
  error[i] <- 1 - mean(knn.pred == test.Y)
}
#plot(nbvoisin,error,xlab="Nombre de voisin pris pour apprentissage",ylab="Erreur estimée")
which.min(error)
```

```
## [1] 6
```

```
print(error[6])
```

```
## [1] 0.1851852
```

According to the result, we obtain the minimum error of **0.1851852** when  $k = 6$ . The K nearest neighbors method generally has poor results in large dimension since the neighbors are actually very far from each other. But here I have done a dimension reduction, so it has a not bad results and no overfitting.

## 1.3 Discriminant Analysis:LDA QDA RDA naive Bayesian

### 1.3.1 LDA

```
lda_data<- lda(y~.,data=train.FDA)
pred<-predict(lda_data,newdata=test.FDA)
table<-table(test.FDA$y,pred$class)
#table
error<-1-sum(diag(table))/nbTst
error
```

```
## [1] 0.1481481
```

```
library(caret)
library(MASS)
folds <- createFolds(y, k = 10, list = TRUE, returnTrain = FALSE)
error_rates_lda <- matrix(nrow = 10, ncol = 1)
for (i in 1:10) {
  train_df_lda <- as.data.frame(X[-folds[[i]],])
  test_df_lda <- as.data.frame(X[folds[[i]],])
  lda_data<- lda(y~.,data=train.FDA)
  lda.pred <- predict(lda_data, newdata = test.FDA)
  error_rates_lda[i, ] <- length(which(as.vector(lda.pred$class) != as.vector(test.FDA$y)))/length(as.v
}
print(mean(as.vector(error_rates_lda)))
```

```
## [1] 0.1481481
```

We get a test error rate with a cross-validation 10-folds is **0.1481481**, the same result with LDA without CV. Linear boundaries therefore give a satisfactory result.

### 1.3.2 QDA

```
qda_data<- qda(y~.,data=train.FDA)
pred<-predict(qda_data,newdata=test.FDA)
table<-table(test.FDA$y,pred$class)
#table
error<-1-sum(diag(table))/nbTst
error
```

```
## [1] 0.4074074
```

We get a very high test error rate: **0.4074074**. This classifier is probably too flexible for our data, so it should be excluded.

### 1.3.4 Naive Bayesian Classifier

```
library(klaR)
```

```
## Warning: package 'klaR' was built under R version 3.3.3
```

```
naivB_data<-NaiveBayes(y~.,data=train.FDA)
pred<-predict(naivB_data,newdata=test.FDA)
table<-table(test.FDA$y,pred$class)
# table
error<-1-sum(diag(table))/nbTst
error
```

```
## [1] 0.2592593
```

We get an error rate of **0.2592593** is therefore worse than that obtained with the LDA (0.14).

### 1.3.5 Regulated discriminant analysis

```
rda_data <- rda(y~.,data=train.FDA, crossval = TRUE)
pred<-predict(rda_data,newdata=test.FDA)
table<-table(test.FDA$y,pred$class)
#table
error<-1-sum(diag(table))/nbTst
error
```

```
## [1] 0.1851852
```

The error rate is **0.1851852**.

## 1.4 Mixture Discriminant Analysis

```
library(mclust)
ind_y = 6
MclustDa_data <- MclustDA(train.FDA[,1:ind_y-1],train.FDA[,ind_y])
#general covariance structure selected by BIC
summary(MclustDa_data, newdata = test.FDA[,1:ind_y-1], newclass = test.FDA[,ind_y])
```

Despite a zero learning error(overfitting), the test error is extremely high (0.4074074).

## 1.5 Tree

### 1.5.1 Decision tree

```
library(tree)
#Full tree
tree_data = tree(as.factor(y)~., train.FDA)
plot(tree_data)
text(tree_data, pretty = 0)
#Cross validation
size<-cv.tree(tree_data)$size
DEV<-rep(0, length(size))
for (i in (1:10))
{
  cv_data = cv.tree(tree_data)
  DEV<-DEV+cv_data$dev
}
DEV <- DEV/10
plot(cv_data$size, DEV, type = 'b')
#Pruning
prune_data = prune.tree(tree_data, best = 7)
plot(prune_data)
text(prune_data, pretty = 0)
#Test Error
y_pred = predict(prune_data, newdata = test.FDA, type = 'class')
table<-table(y_pred, test.FDA$y)
#table
error<-1-sum(diag(table))/nbTst
error
```

The error obtained is **0.4444444**. It is not a good result. We will try to improve it by using Bagging and Random Forest methods.

### 1.5.2 Bagging

```
library(randomForest)
#m = p = 5
bag_data = randomForest(y~., data=train.FDA, mtry=5)
ypred = predict(bag_data, newdata=test.FDA, type = 'response')
table<-table(ypred, test.FDA$y)
error<-1-sum(diag(table))/nbTst
error
```

We have the error rate **0.3333333**.

### 1.5.3 Random Forest

```
#m = sqrt(p)
rdForest_data = randomForest(y~., data=train.FDA,mtry=3)
ypred = predict(rdForest_data, newdata=test.FDA, type = 'response')
table<-table(ypred, test.FDA$y)
```

```
error<-1-sum(diag(table))/nbTst
error
```

We have the error rate **0.2962963**. The errors obtained by bagging and random forest are lower than that obtained initially by the tree pruned.

## 1.6 Support Vector Machine

```
library(e1071)
tune.out = tune(svm, y~., data = train.FDA, kernel = "linear",
               range = list(cost=c(0.01, 0.1, 1, 10, 100), gamma = c(0.1, 1, 10)))
summary(tune.out)
svm_data<-svm(y~., data = train.FDA, kernel = "linear", gamma = 0.1, cost = 1)
ypred = predict(svm_data, newdata=test.FDA)
table<-table(ypred, test.FDA$y)
#table
error<-1-sum(diag(table))/nbTst
error
```

If we use the kernel = **linear**, we obtain the error **0.1851852**, and with the kernel = **radial**, we have **0.2222222**.

## 1.7 Neural Networks (NN)

```
library('nnet')
train.FDA$y = factor(train.FDA$y)
test.FDA$y = factor(test.FDA$y)
model.nnet = nnet(y ~ ., data=train.FDA, size=2, MaxNWts = 20000)
model.nnet.predicted = predict(model.nnet, test.FDA, type="class")
table<-table(model.nnet.predicted, test.FDA$y)
error<-1-sum(diag(table))/nbTst
error
# perfMeasure(model.nnet.predicted, test.FDA$y)
```

The error here is too large and we won't take it.  
The classifier **LDA** is therefore the best for "expressions".

## 2 Données characters\_\_train.txt

```
data2 = data_characters
dim(data2)
set.seed(1)
```

Since I have a large number of observations (10000) and a relatively small number of variables (17), I don't think it necessary to carry out a size reduction. I simply share the data into a test set and a training set.

```
#Generate train data and test data
N<-nrow(data2)
nTrain=floor(3/4*N)
nTest=N-nTrain
train.num<-sample(1:N, nTrain)
```

```
train<-data2[train.num,]
test<-data2[-train.num,]
```

## 2.1 K nearest neighbors(KNN)

```
knn.pred=knn(train.X,test.X,train.Y,k=i)
table(knn.pred,test.Y)
```

I obtain, by cross validation, a minimal error of 0.0760. However, it corresponds to only one neighbor. and I fall into **overfitting**. Here I take the  $k = \sqrt{N}$ .

## 2.2 Analyze Discriminant

### 2.2.1 LDA

```
lda_data<- lda(Y~.,data=train)
pred<-predict(lda_data,newdata=test)
```

I get an error rate of **0.3056**. A classifier with a linear boundary does not seem appropriate.

### 2.2.2 QDA

```
qda_data<- qda(Y~.,data=train)
pred<-predict(qda_data,newdata=test)
```

I obtain an error rate of **0.1224**. The QDA gives good results. This is predictable since I have a **large number of observations**.

### 2.2.3 Naive Bayesian Classifier

```
naivB_data<-NaiveBayes(Y~.,data=train)
pred<-predict(naivB_data,newdata=test)
```

I get an error rate of **0.3636**. The latter is therefore worse than the one obtained with the QDA (0.1304).

### 2.2.4 Regulated discriminant analysis

```
rda_data <- rda(Y~.,data=train, crossval = TRUE)
pred<-predict(rda_data,newdata=test)
```

We have a test error rate **0.122** which is relatively close to the one obtained with QDA. We can thus think that the regularization favored a matrix of covariance peculiar to each class.

## 2.3 Mixture Discriminant Analysis



```
MclustDa_data <- MclustDA(train[, (indice_y+1):ncol(train)], train[, indice_y])
#general covariance structure selected by BIC
summary(MclustDa_data, newdata = test[, (indice_y+1):ncol(train)], newclass = test[, indice_y])
```

I get good results here training error = **0.02666667** and test error = **0.0828**. This classifier is therefore better than that obtained by QDA(0.1304). I can therefore deduce that the distribution of data within classes is closer to a **mixture of Gaussiennes** than to one.

## 2.5 Tree

### 2.5.1 Decision tree

```
library(tree)
tree_data2 = tree(as.factor(Y)~., train)
size<-cv.tree(tree_data2)$size
DEV<-rep(0, length(size))
for (i in (1:10))
{
  cv_data2 = cv.tree(tree_data2)
  DEV<-DEV+cv_data2$dev
}
DEV <- DEV/10
plot(cv_data2$size, DEV, type = 'b')
prune_data2 = prune.tree(tree_data2, best = 17)
y_pred = predict(prune_data2, newdata = test, type = 'class')
table<-table(y_pred, test$Y)
error<-1-sum(diag(table))/nTest
error
```

I use the cross-validation and I get the smallest error for an untagged tree: so I have probably over-learning. This is confirmed by a high test error: **0.616**. This error is therefore not acceptable.

### 2.5.2 Bagging

```
bag_data2 = randomForest(Y~., data = train, mtry = 17)
ypred = predict(bag_data2, newdata = test, type = 'response')
table<-table(ypred, test$Y)
error<-1-sum(diag(table))/nTest
error
```

I obtain the **0.0844** test error. I use  $mtry = p = 5$  by using bagging. It is therefore much better than that obtained previously.

### 2.5.3 Random Forest

```
rdForest_data = randomForest(Y~., data = train, mtry = 4)
ypred = predict(rdForest_data, newdata = test, type = 'response')
table<-table(ypred, test$Y)
error<-1-sum(diag(table))/nTest
error
```

The test error is **0.0532**. It is the best result so far (with the minimum test error).

## 2.6 Neural Networks(NN)

```
model.nnet = nnet(train.Y ~ ., data=train.X, size=2, MaxNWts = 20000)
model.nnet.predicted = predict(model.nnet, test.X, type="class")
table<-table(model.nnet.predicted, test.Y)
error<-1-sum(diag(table))/nTest
error
# perfMeasure(model.nnet.predicted, test.FDA$y)
```

The error is 0.19541.

For the dataset “characters” the best classifier is therefore **Random Forest**.

## 3 Données parole\_train.txt

The dataset has 257 variables. It seems that we need to do a dimension reduction. To be sure, I compared the test errors of an estimated model with and without FDA. The error with FDA is significantly smaller than the error without FDA, so I choose to apply the FDA for all the methods.

```
data3<-data_parole
set.seed(1000)
#Generate train data and test data
N<-nrow(data3);nTrain=floor(3/4*N);nTst=N-nTrain;trainIdx<-sample(1:N, nTrain)
train<-data3[trainIdx,];test<-data3[-trainIdx,]
#FDA (with train data)
lda_data<- lda(y~.,data=train);S<-lda_data$scaling
X<-as.matrix(train[,1:ncol(train)-1]);Z<-X%%S
Z<-as.data.frame(Z);y<-train$y;train_FDA<-cbind(Z,y)
#Apply FDA on test data
X<-as.matrix(test[,1:ncol(test)-1])
Z<-X%%S;Z<-as.data.frame(Z)
y<-test$y;test_FDA<-cbind(Z,y)
```

### 3.1 K nearest neighbors(KNN)

```
knn.pred=knn(train.X,test.X,train.Y,k=i)
table(knn.pred,test.Y)
```

I obtain the minimum error of *0.07460036* when  $k = 17$ .

### 3.2 Discriminant Analysis

#### 3.2.1 LDA

```
lda_data3<- lda(y~.,data=train_FDA)
pred<-predict(lda_data3,newdata=test_FDA)
```

I get an error rate of **0.07282416**. Linear boundaries therefore gives a good result.

### 3.2.2 QDA

```
qda_data3<- qda(y~.,data=train_FDA)
pred<-predict(qda_data3,newdata=test_FDA)
```

I get an error rate of **0.07104796**.

### 3.2.3 Naive Bayesian Classifier

```
NB_data3<-NaiveBayes(y~.,data=train_FDA)
pred<-predict(NB_data3,newdata=test_FDA)
```

The error rate is **0.06749556**.

### 3.2.4 Regulated discriminant analysis

```
rda_data3 <- rda(y~.,data=train_FDA, crossval = TRUE)
pred<-predict(rda_data3,newdata=test_FDA)
```

The error rate is **0.06749556**, it's the same with that of Naive Bayesian Classifier, but I don't have the same table. The different discriminant analysis methods thus give similar test errors. These approaches seem to work well on this data.

## 3.3 Mixture Discriminant Analysis

```
MclustDa_data3 <- MclustDA(train_FDA[,1:indice_y-1],train_FDA[,indice_y])
#general covariance structure selected by BIC
summary(MclustDa_data3, newdata = test_FDA[,1:indice_y-1], newclass = test_FDA[,indice_y])
```

The training error is **0.04682869**, test error is **0.07815275**. This model also gives good result.

## 3.4 Tree

### 3.4.1 Decision tree

```
tree_data3 = tree(as.factor(y)~., train_FDA)
#Pruning
prune_data3 = prune.tree(tree_data3, best = 6)
y_pred = predict(prune_data3, newdata = test_FDA, type = 'class')
table<-table(y_pred, test_FDA$y)
```

The resulting test error is **0.07992895**. It is rather small and is obtained without pruning for 6 leaves. This is probably related to the fact that a dimension reduction has been performed, which limits the complexity of the initial tree.

### 3.4.2 Bagging

```
bag_data = randomForest(y~., data=train_FDA, mtry = 4)
ypred = predict(bag_data, newdata=test_FDA, type = 'response')
```

The error is **0.08703375**.

### 3.4.3 Random Forest

```
rdForest_data = randomForest(y~., data=train_FDA, mtry = 2)
ypred = predict(rdForest_data, newdata=test_FDA, type = 'response')
```

I obtain the error **0.08880995**. Here I take  $mtry = \sqrt{p}$ .

## 3.5 Support Vector Machine

```
tune.out = tune(svm, y~., data = train_FDA, kernel = "radial",
               range = list(cost=c(0.01, 0.1, 1, 10, 100), gamma = c(0.1, 1, 10)))
svm_data<-svm(y~., data = train_FDA, kernel = "radial", gamma = 0.1, cost = 1)
ypred = predict(svm_data, newdata=test_FDA)
```

I have the error **0.06927176** with a cost 0.1 and a gamma 0.1.

## 3.6 Neural Networks(NN)

```
model.nnet = nnet(y ~ ., data=train_FDA, size=2, MaxNWts = 20000)
model.nnet.predicted = predict(model.nnet, test_FDA, type="class")
```

Here we have the error 0.1740675.

The **naïve Bayesian** classifier is therefore the best for the dataset “word”. Besides, I found that all methods have performed well on this dataset.

## Conclusion

- dataset expression : LDA
- dataset character : Random Forest
- dataset parole : Naive Bayesian